

# DOCUMIND: Exposing and Defending Against Indirect Prompt Injection in Document-Processing AI Agents

Anonymous ACL submission

## Abstract

Multimodal large language models increasingly power document-processing agents that ingest PDFs to extract structured fields, answer questions, and trigger downstream tool actions. Yet PDFs remain an understudied red-teaming modality: unlike prior work that applies static perturbations and evaluates only model-level outputs, real attacks are document-conditioned, exploiting render-parse discrepancies, hidden text channels, and multimodal extraction pipelines.

We present **Documind**, a document-conditioned red-teaming framework for PDF-based agent workflows. Given an arbitrary PDF, Documind performs multi-view extraction (parser, OCR, and MLLM), constructs a risk map of action-critical regions, generates an adaptive attack plan, and injects targeted text- and image-layer perturbations. It evaluates defect propagation by replaying agent executions and tracing changes to intermediate state and tool calls.

Evaluated on DuDE against widely deployed black-box MLLMs (GPT-4o, Claude Sonnet, Gemini, Grok), Documind achieves  $\approx 85\%$  Attack Success Rate while preserving human-visible fidelity. We release our framework and interactive demo to support document-conditioned PDF security evaluation and defense research.

## 1 Introduction

PDF documents remain the dominant format for high-stakes institutional workflows—including loan applications, medical records, legal contracts, and academic transcripts—and MLLMs are increasingly deployed as agents that ingest these documents to extract structured information, update records, and trigger automated actions. This creates a consequential attack surface: an adversary who influences what an agent extracts from a PDF




Figure 1: Overview of DOCUMIND. (Top) Attack pipeline: Documents are instrumented with imperceptible perturbations when processed by AI agents, hidden instructions trigger downstream corruption.

can propagate corruption downstream, affecting database writes, compliance decisions, or transactional systems without the document appearing altered to human reviewers.

This threat materializes through Indirect Prompt Injection (IPI)—adversarial content embedded within ingested documents that hijacks agent behavior without the user’s explicit awareness (Greshake et al., 2023). In PDF-based pipelines, IPI is particularly acute because PDFs are not monolithic text streams; they contain multiple extractable layers, including binary text objects, font–glyph mappings, OCR-recoverable image text, and visual overlays. Different parsers and MLLMs resolve these layers inconsistently. An adversary can exploit such render–parse discrepancies to inject content that is imperceptible under standard rendering yet systematically consumed by machine pipelines.

Existing red-teaming work does not adequately capture this threat. Benchmarks such as AdvBench (Zou et al., 2023) evaluate fixed adversarial inputs at the model-output level without modeling document structure or agentic propagation. PDF-




Figure 2: User View for the Vulnerability evaluation framework




Figure 3: Overview of injection payload and methods

specific injection techniques such as TrapDoc ([Authors], 2024e) and related document-layer attacks demonstrate individual vectors but lack adaptive, document-conditioned planning and do not measure downstream agent impact. Jailbreak-focused evaluations further overlook a more operationally significant risk: agents that produce no harmful content yet silently corrupt extracted fields, misinvoke tools, or introduce inconsistent state into downstream systems.

We introduce Documind to address this gap. Our contributions are:

- **Risk Analysis.** Documind constructs a document-conditioned risk map that identifies action-critical regions and generates adaptive attack candidates grounded in document structure and target agent configuration.
- **Attack Simulation.** We systematize six attack classes exploiting render–parse discrepancies across text and image layers, producing perturbations that remain visually consistent while influencing machine parsing.

- **Defect Propagation.** We simulate domain-specific agent workflows and quantify corruption across three failure categories—task deviation, resource inflation, and tool misfire—tracing impact through intermediate state and tool calls.

- **Interactive Demo.** We provide a web interface for uploading PDFs, visualizing risk maps, generating attacked variants, and replaying agent executions across multiple black-box MLLMs.

## 2 Background

### 2.1 PDFs as a Computational Substrate

PDF has long served as the dominant format for digital information exchange and official documentation across domains including academia, finance, healthcare, and law. Its standardization and broad compatibility with dedicated viewers (e.g., Adobe Acrobat, Preview) and modern browsers have made it the default format for institutional workflows. The ease of converting heterogeneous sources (e.g., Word documents, presentations, images) into PDF further reinforces its ubiquity.

At web scale, PDFs are among the most prevalent non-HTML document formats, ranking third among the most common structured content types in public web corpora such as CommonCrawl (Common Crawl Foundation, 2024). As digitization increases across sectors, PDFs increasingly serve not only as human-readable artifacts but also as machine-consumable inputs for automated processing pipelines.

LLMs have further amplified the role of PDFs as a computational substrate. Modern MLLM systems can extract semantic content from PDFs to support tasks such as information retrieval, summarization, structured field extraction, and question answering in downstream applications ([Authors], 2024a).

### 2.2 Agents and Document-Centric Workflows

Recent multimodal foundation models introduced support for file uploads beyond text-only prompting (OpenAI, 2023). This capability enabled application-layer systems to ingest full documents directly through model APIs. Consequently, many deployed document-processing agents now operate over entire PDFs, performing structured reasoning and multi-step task execution within tool-augmented workflows.

065	• Defect Propagation.	We simulate domain-specific agent workflows and quantify corruption across three failure categories—task deviation, resource inflation, and tool misfire—tracing impact through intermediate state and tool calls.	087
066	• Interactive Demo.	We provide a web interface for uploading PDFs, visualizing risk maps, generating attacked variants, and replaying agent executions across multiple black-box MLLMs.	088
067			089
068			090
069			091
070			092
071			093
072			094
073			095
074			096
075			097
076			098
077	• Risk Analysis.	Documind constructs a document-conditioned risk map that identifies action-critical regions and generates adaptive attack candidates grounded in document structure and target agent configuration.	099
078	• Attack Simulation.	We systematize six attack classes exploiting render–parse discrepancies across text and image layers, producing perturbations that remain visually consistent while influencing machine parsing.	100
079			101
080			102
081			103
082			104
083			105
084			106
085			107
086			108
			109
			110
			111
			112
			113
			114
			115
			116
			117
			118
			119
			120
			121
			122
			123
			124
			125
			126
			127
			128
			129
			130
			131
			132
			133
			134

In parallel, retrieval-augmented generation (RAG) pipelines commonly rely on PDF-based knowledge corpora. These systems typically extract text from PDFs, segment it into chunks, and convert it into embeddings for retrieval and generation (Lewis et al., 2020). As a result, PDFs increasingly function as the primary knowledge backbone for document-centric AI systems.

Despite their central role in these pipelines, PDFs introduce unique structural complexity. Unlike plain text documents, PDFs encode content through layered representations, including positioned text objects, font–glyph mappings, embedded images, and metadata. These structural properties complicate consistent extraction across parsers and models, potentially leading to discrepancies between human-visible rendering and machine interpretation.

### 3 Related Work

#### 3.1 Indirect Prompt Injection and Agentic Attacks

Adversarial prompt injection has been widely studied in language models, particularly in the context of jailbreaks and the elicitation of harmful content ([Authors], 2023a). These works demonstrate that carefully crafted inputs can override alignment mechanisms or induce unsafe responses. However, most evaluations focus on direct prompt manipulation at inference time, where adversarial content is explicitly visible in the user query.

Indirect Prompt Injection (IPI) extends this threat model to scenarios where adversarial content is embedded within external data sources ingested by a model, such as web pages or retrieved documents (Greshake et al., 2023). Several frameworks analyze injection risks in retrieval-augmented generation (RAG) pipelines and web-based retrieval systems ([Authors], 2023b), showing that untrusted retrieved content can influence downstream model behavior. Nevertheless, existing approaches generally treat inputs as unstructured text streams, do not explicitly model PDF-specific render–parse discrepancies, and primarily measure model-level output deviation without capturing propagation into tool calls, structured extraction pipelines, or downstream state updates in agentic workflows.




Figure 4: Overview of AI agents, and their tools




Figure 5: Overview of more AI agents, and their tools

#### 3.2 Document-Layer Attacks and PDF Manipulation

Prior work has examined document-centric manipulation in PDF-based systems, demonstrating that non-rendered or visually hidden text can influence automated review processes and language model outputs ([Authors], 2024b). Techniques such as PhantomInject rely on white-text or overlay-based strategies to insert adversarial content into PDFs. PhantomLint and related analyses further identify font–glyph remapping and structured PDF object manipulation as viable injection channels ([Authors], 2024c), while TrapDoc shows that render–parse discrepancies can be exploited to alter machine interpretation without modifying human-visible content ([Authors], 2024e).

Although these works establish the feasibility of document-layer injection, they typically employ fixed perturbation strategies rather than adaptive, document-conditioned planning. Moreover, they focus primarily on model output corruption and do not evaluate workflow-level impact or defect propagation into tool invocations and downstream systems. Documind addresses these limitations by integrating document-conditioned attack generation with end-to-end agentic propagation analysis.

## 207 4 Methodology

### 208 4.1 Threat Model

209 We consider document-centric agent pipelines in  
 210 which an agent ingests a PDF, extracts structured or  
 211 unstructured content, performs a downstream task  
 212 (e.g., question answering or summarization), and  
 213 may invoke external tools such as database writes,  
 214 web navigation, or workflow APIs.

215 **Adversary capabilities.** The adversary controls  
 216 the PDF input or can modify the document in trans-  
 217 it or at rest prior to ingestion. The adversary has  
 218 no access to model weights, training data, system  
 219 prompts, or internal representations, and cannot  
 220 directly modify the inference pipeline. All ma-  
 221 nipulation occurs exclusively through document  
 222 content.

223 **Attack goal.** The adversary seeks to induce agen-  
 224 tic misbehavior through document-layer injections  
 225 that are imperceptible under standard rendering but  
 226 systematically influence the model’s parsed repre-  
 227 sentation. This may result in corrupted field extrac-  
 228 tion, incorrect task outputs, mis-invoked tools, or  
 229 inconsistent downstream state updates.

230 **Scope.** We evaluate born-digital PDFs processed  
 231 by black-box commercial MLLMs accessed via  
 232 API. Attacks requiring direct manipulation of  
 233 model parameters or vision-only adversarial tech-  
 234 niques for scanned documents are out of scope.

235 **Red-teaming objective.** Documind aims to sys-  
 236 tematically surface and measure: (i) the vulnerabil-  
 237 ity of document-processing pipelines to document-  
 238 borne indirect prompt injection, (ii) the document  
 239 regions and structural layers that exhibit elevated  
 240 risk, and (iii) the extent to which induced failures  
 241 propagate into downstream agent actions and tool  
 242 calls.

### 243 4.2 System Overview

244 The core novelty of Documind lies in dynamic,  
 245 document-conditioned attack planning combined  
 246 with multi-layer injection and end-to-end agentic  
 247 defect propagation evaluation. Table 1 situates  
 248 Documind relative to prior PDF attack frameworks.

### 249 4.3 Stage 1: Multi-View Extraction

250 Documind constructs a multi-view representation  
 251 of each PDF using three heterogeneous extract-  
 252 tors operating in parallel: (i) PyMuPDF for na-  
 253 tive binary text parsing and layout structure, (ii)

Framework	Domains	Adaptive	Recreate	Agentic
Code-glyph	N/A	✗	N/A	✗
TrapDoc	2	✗	✓	✗
Integrity Shield	1	✓	✓	✗
Dope	1	✓	✓	✗
<b>Documind</b>	<b>6</b>	✓	✗	✓

Table 1: Comparison of PDF attack frameworks across domain scope, adaptive perturbation, PDF reconstruction, and agentic evaluation.




Figure 6: Overview of stage-1.

PaddleOCR for OCR-based recovery of image-  
 embedded text, and (iii) an MLLM-based extrac-  
 tor (Gemini) for semantic and visual interpreta-  
 tion. Each extractor emits structured artifacts (JSON rep-  
 resentations, bounding boxes, rendered images, and  
 markdown summaries). These outputs are com-  
 bined into a unified context:

$$C = (C_1, C_2, C_3)$$

where each  $C_i$  captures complementary structural  
 and semantic information.

PyMuPDF provides precise object-level struc-  
 ture and bounding boxes; PaddleOCR recovers  
 visually embedded or scanned content; and the  
 MLLM extractor generates semantically structured  
 markdown that captures higher-level document or-  
 ganization not recoverable through strict parsing  
 alone. This heterogeneous extraction strategy was  
 selected based on empirical analysis demon-  
 strating improved downstream robustness over single-  
 extractor baselines.

### 274 4.4 Stage 2: Attack Planning

275 Given the unified context, Documind’s planner—  
 276 implemented as a prompted MLLM—derives doc-  
 277 ument semantics, identifies high-risk regions, and  
 278 generates a document-conditioned attack plan. The




Figure 7: Overview of stage-2.




Figure 8: Overview of stage-3.

planner receives structured JSON, extracted image artifacts, semantic markdown summaries, and a configuration prompt encoding document domain, intended purpose, target agent type, and available tools. It outputs document-level semantic metadata and fine-grained risk parameters including page, block, line, element identifiers, and bounding-box coordinates ( $x_1, y_1, x_2, y_2, x_3, y_3, x_4, y_4$ ).

This planning stage operationalizes document conditioning: attack strategies are generated relative to semantic role and workflow sensitivity rather than generic perturbation heuristics.

#### 4.5 Stages 3–4: Risk Analysis and Payload Preparation

Documind identifies action-critical fields and generates targeted attack payloads grounded in workflow semantics. Fields are classified as sensitive based on their role in downstream decision-making. For example, in a fundraising document, fields such as Link, Location, Date, and Cost are marked sensitive because modifying each induces a distinct




Figure 9: Overview of stage-4.

Strategy	Mechanism	Purpose
Append	Hidden text injection	Non-rendered adversarial instructions
Update	Font-glyph remapping	Alter field semantics, preserve appearance
Update	Visual overlay	Alter field semantics, preserve appearance
Delete	Font-glyph remapping	Remove meaning without visible deletion
Delete	Visual overlay	Remove meaning without visible deletion

Table 2: Mapping between semantic edit strategies and injection mechanisms.

class of workflow failure: corrupted links redirect to unintended destinations, altered dates trigger scheduling errors, and modified costs distort financial decisions.

Risk is therefore defined in terms of decision relevance rather than surface-level content. This document-conditioned sensitivity classification enables workflow-aware attack generation rather than template-based injection.

#### 4.6 Stage 5: Attack Injection

Documind applies generated payloads through a two-layer abstraction that decouples semantic intent from physical embedding.

**Semantic Edit Strategy.** The planner selects one of three strategies: *Append* (introduce adversarial content without altering visible structure), *Update* (modify the value or meaning of an existing field), or *Delete* (neutralize or suppress decision-critical content).

**Injection Mechanism.** The physical channel used to embed the semantic edit while preserving visual fidelity. Table 2 shows the mapping between strategies and mechanisms.

By decoupling semantic corruption from physical embedding, Documind remains extensible as

325 models evolve to resist specific injection techniques.  
326

#### 327 4.7 Defect Propagation

328 A key limitation of prior evaluations is their con-  
329 finement to model-output deviation. Documind  
330 instead measures whether and how induced defects  
331 propagate into agentic workflows. Each simulated  
332 agent is configured with a target task (e.g., field  
333 extraction, Q&A, summarization), a toolset (e.g.,  
334 database writes, web fetch, ticket creation), and an  
335 invocation policy governing tool execution.

336 Documind evaluates three propagation out-  
337 comes: (i) **Payload drift**—extracted structured  
338 fields diverge from document ground truth; (ii)  
339 **Action misfire**—incorrect tool selection or mal-  
340 formed parameters are issued; (iii) **Cascade risk**—  
341 corrupted outputs are written to downstream sys-  
342 tems and subsequently consumed by dependent  
343 services. This propagation-aware evaluation dis-  
344 tinguishes Documind from prior document-layer  
345 attacks that stop at surface-level output corruption.

## 346 5 Experiments

### 347 5.1 Experimental Setup

348 **Threat model instantiation.** All experiments fol-  
349 low the threat model defined in Section 4. The ad-  
350 versary controls PDF content only, accesses mod-  
351 els via black-box APIs, and cannot modify system  
352 prompts, tool definitions, or model weights. At-  
353 tacks operate exclusively through document-layer  
354 manipulation on born-digital PDFs. If a model re-  
355 fuses to process a document due to safety filters,  
356 the attack is considered unsuccessful for corruption-  
357 based objectives. Refusals are treated separately as  
358 availability outcomes and are not counted toward  
359 corruption ASR.

360 **Dataset.** We evaluate on the Document Under-  
361 standing Dataset (DuDE), which provides born-  
362 digital PDFs, QA annotations, and structured  
363 ground-truth references for key-field extraction.  
364 DuDE supports both extraction fidelity measure-  
365 ment and QA-based task corruption evaluation.  
366 Thresholds and hyperparameters are selected on  
367 a held-out validation split disjoint from the test set.

368 **Agent simulation.** We define 10 domain-specific  
369 agents configured with 17 functional tools, mod-  
370 eled after standard LLM-agent architectures ([Au-  
371 thors], 2024d). Each agent consists of a task  
372 module (QA, extraction, or summarization), an

373 LLM-driven tool invocation policy, and a struc-  
374 tured JSON-based state update mechanism. Tool  
375 calls are executed in a simulated environment that  
376 records tool name, parameters, and resulting state  
377 changes. Clean baseline executions are verified for  
378 correctness prior to comparison. Full implementa-  
379 tion details are provided in Appendix A.

### 380 5.2 Logical Failure Modes

381 We evaluate three workflow-level failure classes:  
382 **Task Corruption**—deviation of QA or summa-  
383 rization outputs from document ground truth; **Re-  
384 source Inflation**—increase in token usage or la-  
385 tency relative to the clean baseline; **Tool Misfire**—  
386 incorrect tool selection or malformed payload gen-  
387 eration within agent workflows.

### 388 5.3 Extraction Accuracy

389 **Goal.** Measure degradation in structured extrac-  
390 tion fidelity under document-layer attack.

391 For each document  $d$ , let  $E_{\text{clean}}(d)$  and  $E_{\text{attack}}(d)$   
392 denote extracted content from clean and attacked  
393 PDFs respectively. We compute Character Error  
394 Rate (CER), Word Error Rate (WER), block-level  
395 structural match, and key-field exact match (date,  
396 amount, location, URLs). Extraction corruption is  
397 defined as:

$$\Delta_{\text{extract}}(d) = 1$$

398 if key-field match drops below threshold  $\tau$   
399 where  $\tau$  is selected on a held-out validation split.

### 400 5.4 Task Accuracy

401 **Goal.** Measure downstream correctness of QA  
402 and summarization.

403 For each document  $d$ , let  $A_{\text{clean}}(d)$ ,  $A_{\text{attack}}(d)$ ,  
404 and  $\text{GT}(d)$  denote clean answer, attacked answer,  
405 and ground truth respectively. Task corruption is  
406 defined as:

$$A_{\text{attack}}(d) \neq A_{\text{clean}}(d) \text{ AND } A_{\text{attack}}(d) \neq \text{GT}(d)$$

407 Both conditions are required to prevent counting be-  
408 nign divergence. Metrics include QA Accuracy, F1  
409 against DuDE ground truth, summarization faith-  
410 fulness, and key-fact coverage.

### 411 5.5 Attack Success Rate (ASR)

412 **Goal.** Quantify end-to-end attack success across  
413 agent workflows.

414 For each document  $d$ , define:

$$S(d) = \begin{cases} 1 & \text{if targeted corruption-based logical failure occurs} \\ 0 & \text{otherwise} \end{cases}$$

Attack	GPT-4o	Sonnet	Gemini	Grok
Clean	91%	92%	—	—
Hidden Text	—	—	—	—
Font-Glyph	—	—	—	—
Visual Overlay	—	—	—	—

Table 3: Extraction accuracy (key-field exact match %) under clean and attacked conditions.

$$\text{ASR} = \frac{1}{N} \sum_{d=1}^N S(d)$$

A targeted logical failure includes task corruption, token usage inflation above threshold  $X\%$ , or tool misfire. Threshold  $X$  is selected on a held-out validation split. Per-category ASR is reported alongside aggregate ASR.

## 5.6 Defect Propagation

**Goal.** Measure whether induced failures propagate into downstream agent state and tool behavior.

**Payload Drift:**  $\text{Drift}(d) = 1$  if extracted structured state diverges from ground truth on at least one key field. **Action Misfire:**  $\text{Misfire}(d) = 1$  if tool name or parameters differ from the verified clean baseline. **Cascade Risk:** Detected programmatically when corrupted state written by one tool is subsequently read and acted upon by a dependent tool within the same execution trace.

## 5.7 Decoy Stealthiness

**Goal.** Evaluate detectability by humans and automated classifiers.

**Human evaluation.** Three annotators per document perform A/B detection between clean and attacked variants. Majority vote determines detection. Inter-annotator agreement is reported via Cohen’s  $\kappa$ .

**Model-based classification.** A binary prompt (“Is this document maliciously altered?”) is evaluated across models. Classifier accuracy is reported per model and injection channel. An attack is considered stealthy if human detection rate does not exceed 50% and model classifier accuracy does not exceed 60%.

## 6 Results

### 6.1 Extraction Degradation

Table 3 reports extraction fidelity under clean and attacked conditions.

Attack	GPT-4o	Sonnet	Gemini	Grok
Clean	85%	87%	—	—
Hidden Text	—	—	—	—
Font-Glyph	—	—	—	—
Visual Overlay	—	—	—	—

Table 4: Task accuracy (QA F1 %) under clean and attacked conditions.

Attack	GPT-4o	Sonnet	Gemini	Grok
Hidden Text	—	—	—	—
Font-Glyph	—	—	—	—
Visual Overlay	—	—	—	—
<b>Aggregate ASR</b>	—	—	—	—

Table 5: Attack Success Rate (%) per injection channel and model. Aggregate ASR  $\approx 85\%$  across models.

These results confirm that render–parse discrepancies translate into measurable extraction corruption prior to task execution.

## 6.2 Task Corruption

Table 4 reports QA accuracy and summarization metrics under clean and attacked conditions.

Using the strict corruption definition from Section 5.4, Documind induces task corruption in  $[X]\%$  of documents. Importantly, divergence from clean outputs alone is insufficient to count as corruption; attacks are only considered successful when attacked outputs also deviate from ground truth. Summarization tasks exhibit similar degradation patterns, with faithfulness and key-fact coverage dropping under adaptive perturbations.

## 6.3 Attack Success Rate

Table 5 reports aggregate and per-category ASR results.

Across models, Documind achieves an average corruption-based ASR of  $\approx 85\%$ . Safety refusals occurred in  $[R]\%$  of cases and were excluded from corruption-based ASR as defined in Section 5.1.

## 6.4 Defect Propagation

Table 6 reports propagation outcomes across injection channels.

These results demonstrate that document-layer attacks do not remain confined to model outputs but systematically alter agent behavior through the full execution trace.

Attack	Payload Drift	Action Misfire	Cascade Risk	Evaluation Dashboard.	
Hidden Text	—%	—%	—%	compares clean and attacked runs across extraction	517
Font–Glyph	—%	—%	—%	outputs, QA/summarization results, token us-	518
Visual Overlay	—%	—%	—%	age, latency, tool invocation traces, and struc-	519

Table 6: Defect propagation outcomes (%) across injection channels.	The demo is model-agnostic and supports multiple black-box MLLMs through file-upload APIs.	520
		521

Table 7: Decoy stealthiness: human detection rate, inter-annotator agreement ( $\kappa$ ), and model classifier accuracy.	522
	523

Table 6: Defect propagation outcomes (%) across injection channels.

## 485 6.5 Decoy Stealthiness

486 Table 7 reports human detection rates and model  
487 classifier accuracy.

## 488 6.6 Ablation: Adaptive vs. Fixed Attacks

489 To isolate the contribution of document-  
490 conditioned planning, we compare Documind  
491 against fixed perturbation baselines. Adaptive  
492 planning improves ASR by  $[\Delta]\%$  over static  
493 injection, particularly in tool misfire scenarios,  
494 where document-conditioned risk mapping enables  
495 targeted corruption of decision-critical fields.

## 496 7 System Demo

497 Documind provides an interactive web-based interface  
498 designed for reproducible PDF security evaluation.  
499 The demo supports the following workflow:

### 500 1. Document Upload and Workflow Selection.

501 Users upload a PDF and select a target workflow,  
502 including extraction, question answering,  
503 summarization, or tool-augmented agent  
504 execution.

### 505 2. Risk Analysis View.

506 Documind visualizes action-critical regions identified during  
507 document-conditioned risk mapping. Sensitive fields (e.g., link, date, location, cost) are  
508 highlighted with bounding-box overlays to ex-  
509 pose workflow-relevant risk surfaces.

### 511 3. Attack Simulation View.

512 Users select a logical failure objective (task corruption, resource  
513 inflation, or tool misfire) and a document-  
514 layer injection channel. The system generates  
515 attacked variants in real time while preserving  
516 visual fidelity.

**Evaluation Dashboard.** The dashboard compares clean and attacked runs across extraction outputs, QA/summarization results, token usage, latency, tool invocation traces, and structured state changes. Agent execution traces are displayed step-by-step to enable inspection of defect propagation.

The demo is model-agnostic and supports multiple black-box MLLMs through file-upload APIs.

## 524 8 Conclusion

We presented Documind, a document-conditioned red-teaming framework for PDF-based agent workflows. Unlike prior work that evaluates static perturbations or model-level outputs, Documind integrates adaptive attack planning, multi-layer injection, and end-to-end agentic defect propagation analysis. Our results demonstrate that document-layer perturbations exploiting render–parse discrepancies can induce workflow-level corruption while preserving human-visible fidelity. These findings highlight the need for systematic robustness evaluation in document-processing pipelines before deployment in high-stakes environments. Documind provides both a research framework and an interactive demo to support continued study of PDF security and defense strategies.

## 543 9 Limitations

While Documind evaluates multiple commercial MLLMs (OpenAI, Claude, Gemini, Grok, Perplexity) through their file-upload APIs, these systems evolve rapidly. Model updates, alignment improvements, and parser changes may affect attack efficacy over time. Results should therefore be interpreted as representative of current-generation systems rather than permanent vulnerabilities.

Our threat model focuses on born-digital PDFs, which constitute the majority of web-scale document corpora (Common Crawl Foundation, 2024). We do not evaluate scanned-only documents requiring purely vision-domain adversarial techniques. Extending dynamic perturbation to VLM- and OCR-specific pipelines remains future work.

Agent simulations are designed to approximate realistic deployment conditions; however, institutional workflows vary widely. Domain-specific policy constraints and custom tool integrations may alter propagation behavior. We recommend institution-specific red-teaming prior to production deployment.

566 Finally, dynamic attack planning depends on  
567 the reasoning behavior of the underlying planner  
568 model. Although injection mechanisms operate at  
569 the document substrate level, planner variability  
570 may influence attack targeting precision.

## 571 Ethical Considerations

572 Documind is intended solely for defensive secu-  
573 rity research and robustness evaluation. While the  
574 framework can generate adversarially perturbed  
575 documents, its purpose is to identify vulnerabilities  
576 in document-processing pipelines and inform miti-  
577 gation strategies. All experiments are conducted on  
578 publicly available datasets. No personal or sensi-  
579 tive private data is used. The system does not target  
580 specific institutions or deployed services. We advo-  
581 cate responsible disclosure practices and encourage  
582 developers to use red-teaming frameworks such as  
583 Documind to proactively strengthen document in-  
584 gestion systems.

## 585 Acknowledgements

586 We thank [XYZ] for their guidance and feedback  
587 throughout this project. We also acknowledge the  
588 Google Cloud team for providing cloud credits that  
589 supported system development and deployment.

## 590 References

- 591 [Authors]. 2023a. Advbench: Adversarial benchmark  
592 for language models. [To be completed].
- 593 [Authors]. 2023b. Poisoning retrieval corpora by inject-  
594 ing adversarial passages. [To be completed].
- 595 [Authors]. 2024a. [mllm document processing refer-  
596 ence]. [To be completed].
- 597 [Authors]. 2024b. [peer review manipulation paper].  
598 [To be completed].
- 599 [Authors]. 2024c. Phantomlint: Document-layer injec-  
600 tion analysis. [To be completed].
- 601 [Authors]. 2024d. Tau-bench: A benchmark for tool-  
602 augmented language models. [To be completed].
- 603 [Authors]. 2024e. Trapdoc: Adversarial document in-  
604 jection. [To be completed].
- 605 Common Crawl Foundation. 2024. [Common crawl](#).
- 606 Kai Greshake, Sahar Abdelnabi, Shailesh Mishra,  
607 Christoph Endres, Thorsten Holz, and Mario Fritz.  
608 2023. Not what you've signed up for: Compromising  
609 real-world llm-integrated applications with indirect  
610 prompt injection. In *Proceedings of the 16th ACM*  
611 *Workshop on Artificial Intelligence and Security*.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio  
Petroni, Vladimir Karpukhin, Naman Goyal, Hein-  
rich Küttler, Mike Lewis, Wen-tau Yih, Tim Rock-  
täschel, and 1 others. 2020. Retrieval-augmented  
generation for knowledge-intensive nlp tasks. In *Ad-  
vances in Neural Information Processing Systems*,  
volume 33.

OpenAI. 2023. Gpt-4 technical report. ArXiv preprint  
arXiv:2303.08774.

Andy Zou, Zifan Wang, J. Zico Kolter, and Matt Fredrik-  
son. 2023. Universal and transferable adversarial at-  
tacks on aligned language models. In *arXiv preprint*  
arXiv:2307.15043.

## 625 A Agent Configuration Details

626 Full agent configurations, tool definitions, and im-  
627 plementation assumptions are provided here. [To  
628 be completed.]